

What changed with TradingAgents?

A plain-language owner update after about 15 hours of work • July 26, 2026

The honest one-minute answer

The system is much better at researching, checking itself, recovering, and testing new strategies without asking you what to buy. But the improved version is still in its protected workshop, not running the live account. The old live system is correctly frozen because three safety records are stale or out of order.

3 / 5

current phase

≈61%

full goal complete

115

worker assignments

3,160

last full checks passed

Where are we right now?

Current phase: Phase 3 of 5. We are teaching the firm how to create, compare, reject, and safely advance its own strategies.

Overall progress: about 61%, or roughly 23 of 38 major work blocks. This is an estimate, not a fake-precise score.

Capability improvement: about +65% from the version we restarted. It has many more useful abilities and stronger internal checks.

Production confidence: about +42%. That number is lower because the new abilities are not yet installed into the live schedule and have not survived five real market days.

Time remaining: about 12 to 20 more active agent hours, then at least five market days of proof. A realistic finish window is 6 to 10 calendar days after work resumes, if the proof run stays clean.

Money reality

The live Alpaca account currently shows \$196.70 from the original \$200. That is down \$3.30. The work so far has improved the machine, not proven that it can grow the money yet.

What changed?

1. It can recover without waiting for you

We built a clear chain of command for ordinary failures. One role detects the problem, one role repairs it, and a separate role checks the repair before trading can restart. The owner is no longer the normal repair person.

2. It can invent bounded strategy variations

The system can now adjust allowed strategy ingredients, build a new candidate, test it on its own pretend \$200 account, compare the result, preserve the evidence, and prepare the winner for paper trading. It still cannot secretly turn a new idea into a live order.

3. It remembers why decisions were made

Agents now hand each other short, structured decision packets instead of giant loose conversations. The next agent can see the claim, the evidence, the risk, and what must happen next. Work can pause and resume without starting over.

4. Research is harder to fool

Research now rejects old prices, future information accidentally leaking into a past test, broken data providers, and mismatched ticker names. When a source fails, the system records the failure and uses a controlled fallback instead of pretending the answer is complete.

5. Strategy promotion has a paper trail

A candidate cannot advance on a persuasive paragraph alone. Its identity, test window, evidence, decision, and later outcome must match. A second checker reviews the packet before it can move forward.

6. The live side stays fail-closed

If the system cannot prove that every safety check happened in the right order, it freezes. That is why the live account is blocked now. The freeze is inconvenient, but it is the correct behavior for an unknown safety state.

How much work actually ran?

115	44	48	16	7
separate assignments	build or repair	independent reviews	research or design	acceptance tests

Those 115 assignments also received 166 follow-up work or review turns. The last committed full examination passed 3,160 checks, with one intentionally skipped check and no failing check.

Important: a passed check means the machinery behaved as designed under test. It does not mean the strategy will make money in the market.

The five phases

Phase 1: Recovery and control · 8 / 8 · 100% · **DONE**

Who can freeze, repair, verify, and re-arm is now explicit. Ordinary failures no longer need to be routed to you.

Phase 2: Better research and data · ≈8 / 10 · 80% · **MOSTLY DONE**

Stale data, provider failures, symbol mismatches, and future-information leaks are contained. A few broad provider-quality slices remain.

Phase 3: Communication and strategy evolution · ≈5.75 / 9 · 64% · **CURRENT**

Agent handoffs and learning records work. Strategy creation, testing, paper staging, and shadow evidence work. Final promotion and rollback behavior is still being finished.

Phase 4: Put the new firm into production · ≈1.5 / 7 · 20% · **STARTED**

The schedule and role map exist, but the new chain of command has not replaced the old production prompts yet.

Phase 5: Real-market proof and final audit · 0 / 4 · 0% · **NOT STARTED**

The final system must survive five market days, prove order reconciliation, prove safe recovery, and receive a final independent review.

What has been proven to work?

The Alpaca paper and live connections answer read-only checks, and the broker currently reports no open orders.

The schedule contains ten jobs, and every one is set to notify you only when a run fails.

New strategy candidates can be created within strict limits, tested on isolated pretend money, stored with permanent receipts, and staged for paper review.

Research can survive a missing provider, reject stale market data, normalize ticker names, and record what information was available at the time.

The improved branch completed a full 3,160-check examination before the current small prerequisite correction.

The latest overnight run ranked 40 symbols and completed two of three deep company analyses, using a controlled fallback for the third.

What does not work yet?

The improved system is not deployed into the live production schedule.

Live trading remains frozen because safety timing, checker strength, and overnight evidence are not current enough to authorize trading.

The strategy-evolution path is intentionally unable to submit broker orders. Promotion, demotion, rollback, and live authority are later gates.

The visual n8n dashboard and the YouTube transcript helper are down because Docker is not running.

One optional financial-data provider is unavailable because its paid access key is missing. The system falls back, but transcript and earnings detail can be thinner.

Historical forecast scoring is noisy and mostly poor. New scoring and down-weighting controls exist, but they have not earned forward-market trust yet.

The required five-market-day proof has not begun.

Today's production snapshot

Connected, but correctly frozen

Alpaca is answering, the account is active, and there are no known open orders. The freeze comes from our own safety rules, not from a broken broker connection.

Live account value: \$196.70. Buying power: \$86.42.

Known open orders: zero live and zero paper.

Schedule: nine jobs on, one job paused. Overnight Research is the paused job.

Notifications: failed runs only for all ten jobs.

n8n: the safe runner answers, but the visual screen is down because Docker is off.

Who works when? All 10 scheduled jobs

Seven roles, ten scheduled jobs

Some roles run more than one job. All times below are Central Time. Only the Market Supervisor may start the normal live-order path, and even it must pass every safety gate.

1. Wake Controller: Schedule Controller

When: Mon to Fri, 6:45 AM **Status:** ON

Checks whether the market is open, then wakes the market-day jobs. It does not trade.

2. Overnight Research: Research & Strategy Council

When: Mon to Fri, 3:30 AM **Status:** PAUSED

Studies and ranks opportunities before the market opens. It produces research only, never orders.

3. Preopen Validation: Integrity Verifier

When: Mon to Fri, 8:15 AM **Status:** ON

Runs a dry safety rehearsal before trading time. Unknown or stale evidence means no go.

4. Safety Sentinel: Integrity Verifier

When: Hourly, 8:25 AM to 2:25 PM **Status:** ON

Checks the guardrails just before each trading cycle. It can freeze the system but cannot trade.

5. Market Supervisor: Execution Operator

When: Hourly, 8:35 AM to 2:35 PM **Status:** ON

Reviews approved intent and may start the normal order wrapper only after every gate says yes.

6. Execution Board: Portfolio Executive

When: Hourly, 8:50 AM to 2:50 PM **Status:** ON

Judges portfolio risk and execution quality. It may freeze trading, but it cannot place an order or re-arm itself.

7. Paper Tournament: Research & Strategy Council

When: Mon to Fri, 9:05 AM **Status:** ON

Lets candidate strategies compete with pretend money. It cannot promote itself into live trading.

8. Self-Healer: Reliability Controller

When: 7:03, 9:03, 11:03 AM; 1:03, 3:03 PM **Status:** ON

Repairs ordinary local and observer problems. The production version still needs the new independent re-arm chain installed.

9. Daily Report: Reporting Utility

When: Mon to Fri, 3:30 PM **Status:** ON

Summarizes the day and sends the report. It has no trading authority.

10. Sleep Controller: Schedule Controller

When: Mon to Fri, 4:45 PM **Status:** ON

Puts market-day jobs to sleep and decides whether overnight research should run.

What does the system depend on?

Codex scheduled tasks · ACTIVE

Think of it as: Alarm clocks for the workers.

What it does: Wake each job at the right Central Time and keep the daily rhythm moving.

Without it: The trading program could still be started by hand, but it would not operate as a dependable daily firm.

Alpaca · WORKING

Think of it as: The bank counter and order desk.

What it does: Holds the \$200 experiment account, reports cash and positions, and receives an order only after authorization.

Without it: The agents could research and paper-test, but they could not see or trade the real experiment account.

OpenRouter · WORKING

Think of it as: A switchboard that connects us to several thinking models.

What it does: Lets overnight research use DeepSeek and Qwen without wiring each model separately.

Without it: The system can still use simpler local rules and fallbacks, but deep research becomes weaker or more limited.

Yahoo Finance / yfinance · WORKING

Think of it as: A free market-information window.

What it does: Provides prices, news, option context, short interest, and event dates.

Without it: Research would have bigger holes or would need more paid data services.

LangGraph and LangChain · BUILT IN

Think of it as: The assembly-line map.

What it does: Tell the analysts, debaters, risk checker, and manager whose turn it is, what to pass next, and where to resume.

Without it: We would have a pile of separate prompts instead of a coordinated team that can pause and continue.

Local receipts and memory · BUILT IN

Think of it as: A filing cabinet with signed slips.

What it does: Keeps decisions, evidence, checkpoints, test outcomes, and order reconciliation records on this Mac.

Without it: The agents could forget why something was approved, repeat work, or be unable to prove what happened.

Git · BUILT IN

Think of it as: A tamper-evident version stamp.

What it does: Proves exactly which strategy and rules produced each result, so a later review can reproduce it.

Without it: Promotion and rollback evidence would be much easier to mix up or rewrite accidentally.

Email service · CONFIGURED

Think of it as: The mail carrier.

What it does: Sends the daily owner report and failure messages.

Without it: The reports would still be saved locally, but you would need to open the machine to read them.

What is n8n, and do we need it?

Short answer: useful, but not the engine

n8n is a visual control room. It lets us draw a workflow as boxes and arrows, press a safe button, and see what happened. It is not allowed to hold Alpaca secrets or place trades.

What it improves: easier manual inspection, a visual dashboard, repeatable evaluation workflows, and a safer way to call a small allow-listed set of read-only helper actions.

What we have without it: the Codex schedules, research, safety checks, broker connection, and core trading workflow still run. We lose the visual control room and some convenient observer tools, not the core autonomous firm.

Current status: the small n8n runner is alive, but the visual dashboard is down because Docker is off.

Docker explained: Docker is a lunchbox that keeps helper services packaged with everything they need. When the lunchbox is closed, n8n's screen and the YouTube transcript helper cannot start. Alpaca and the core trading system do not depend on it.

Other optional helpers

Financial Modeling Prep is a paid research source. Its key is missing, so the system skips it and uses other sources. We lose some earnings and transcript depth, not basic operation.

MCP and Composio connectors are read-only bridges to outside information or apps. They can save manual work, but they do not receive trading authority.

MiroFish and Zep are older optional simulation and memory experiments. They are not required for the current autonomous-firm plan, and their past scoring was not strong enough to trust as the main brain.

Why is live trading frozen?

The blocker in one sentence

The broker is healthy, but the live system cannot prove that the safety check finished before the trading cycle, the sentinel is set to a weaker thinking level than required, and the overnight research record is stale.

Is your approval still the blocker?

No. Your earlier approval removed the human-authorization issue. The remaining blocker is machine evidence and rollout work. The system should repair ordinary failures by itself once the new recovery chain is installed.

How do we unfreeze it safely?

1. Finish and independently review the small strategy prerequisite correction that was paused for this report.
2. Install the corrected schedule order, stronger sentinel setting, and fresh overnight evidence rules into production.
3. Run a clean read-only broker reconciliation and confirm zero unknown orders or positions.
4. Let the independent verifier issue a short-lived re-arm decision. The repair role cannot approve its own repair.
5. Start with shadow and paper observation, then a tiny canary only after every gate remains green.

What happens next?

Immediate checkpoint: complete the paused prerequisite review, run the long repeat tests and the full examination, get two fresh independent reviews, and save the accepted result.

Next strategy work: finish safe promotion, demotion, rollback, and selection so poor strategies lose influence and strong strategies can earn more paper responsibility.

Production rollout: replace the old automation instructions with the new seven-role chain, repair the overnight schedule, and prove that repair and re-arm are independent.

Final proof: run at least five market days, reconcile every broker action, review every freeze and repair, and compare real outcomes with the system's forecasts.

The bottom line

You were not supposed to choose the trades, and the final system still will not ask you to. We have built most of the internal firm and its safety rails. The unfinished work is connecting that improved firm to production, proving it for five market days, and showing that it can grow the account rather than merely operate safely.

How to read these numbers

The percentages are a rough owner-level estimate based on completed major work blocks, not lines of software. “Capability” means what the system can now do in the protected workshop. “Production confidence” means how much of that behavior is installed, observed, and proven against the real schedule. Market results can still be worse than expected even when the machinery passes every safety check.